


# Creating Cloud Resources with Terraform



Michael Bright , Terraform & Vault Associate



<https://linkedin.com/in/mjbright>



@mjbright



@mjbright

# Webinar Agenda

## 1. Introduction to Terraform

- Configurations, Providers
- Installation, Workflow
- Variables: Input, Output, Local

## 3. Working with Terraform Registry

- Data sources
- Modules


## 2. Control structures

- Replicas using Count

## 4. Terraform in Production

- Provisioners
- Managing State

# 1. Introduction



Terraform Concepts



<https://linkedin.com/in/mjbright>



@mjbright



@mjbright

# HashiCorp Products & Projects



HashiCorp




HashiCorp

**Terraform**

+ Terraform Cloud/  
Terraform Enterprise  
[ Sentinel ]

Terraform CDK  
Python/Typescript



HashiCorp

**Consul**



HashiCorp

**Vault**



HashiCorp

**Packer**



HashiCorp

**Vagrant**



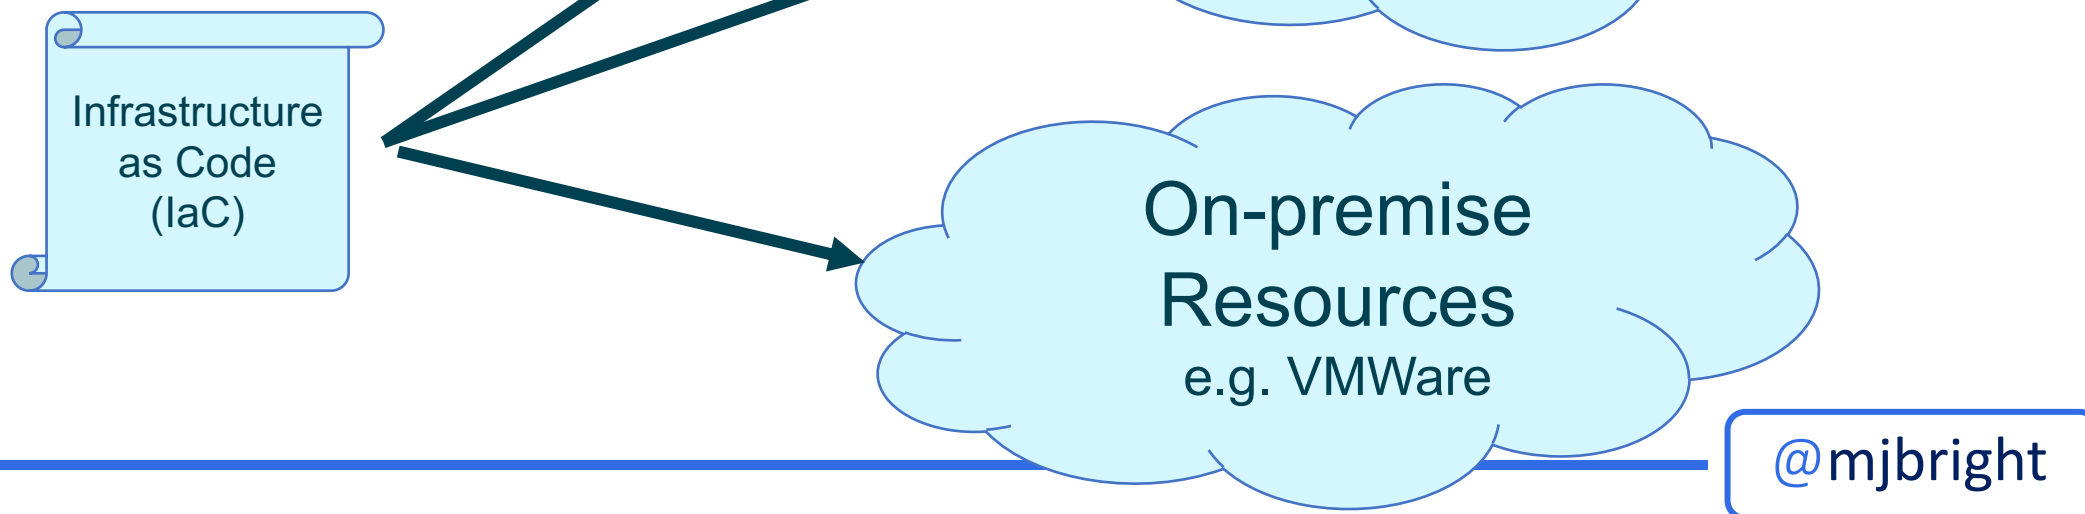
HashiCorp

**Nomad**

# Introduction

“Infrastructure as Code” (IaC) brings

- Automation, repeatability
- Documentation
- Version Control, Audit trail
- Validation, Testability
- Reuse



# Introduction

Infrastructure as Code (IaC) is a descriptive model, where source code describes the infrastructure resources to be created.

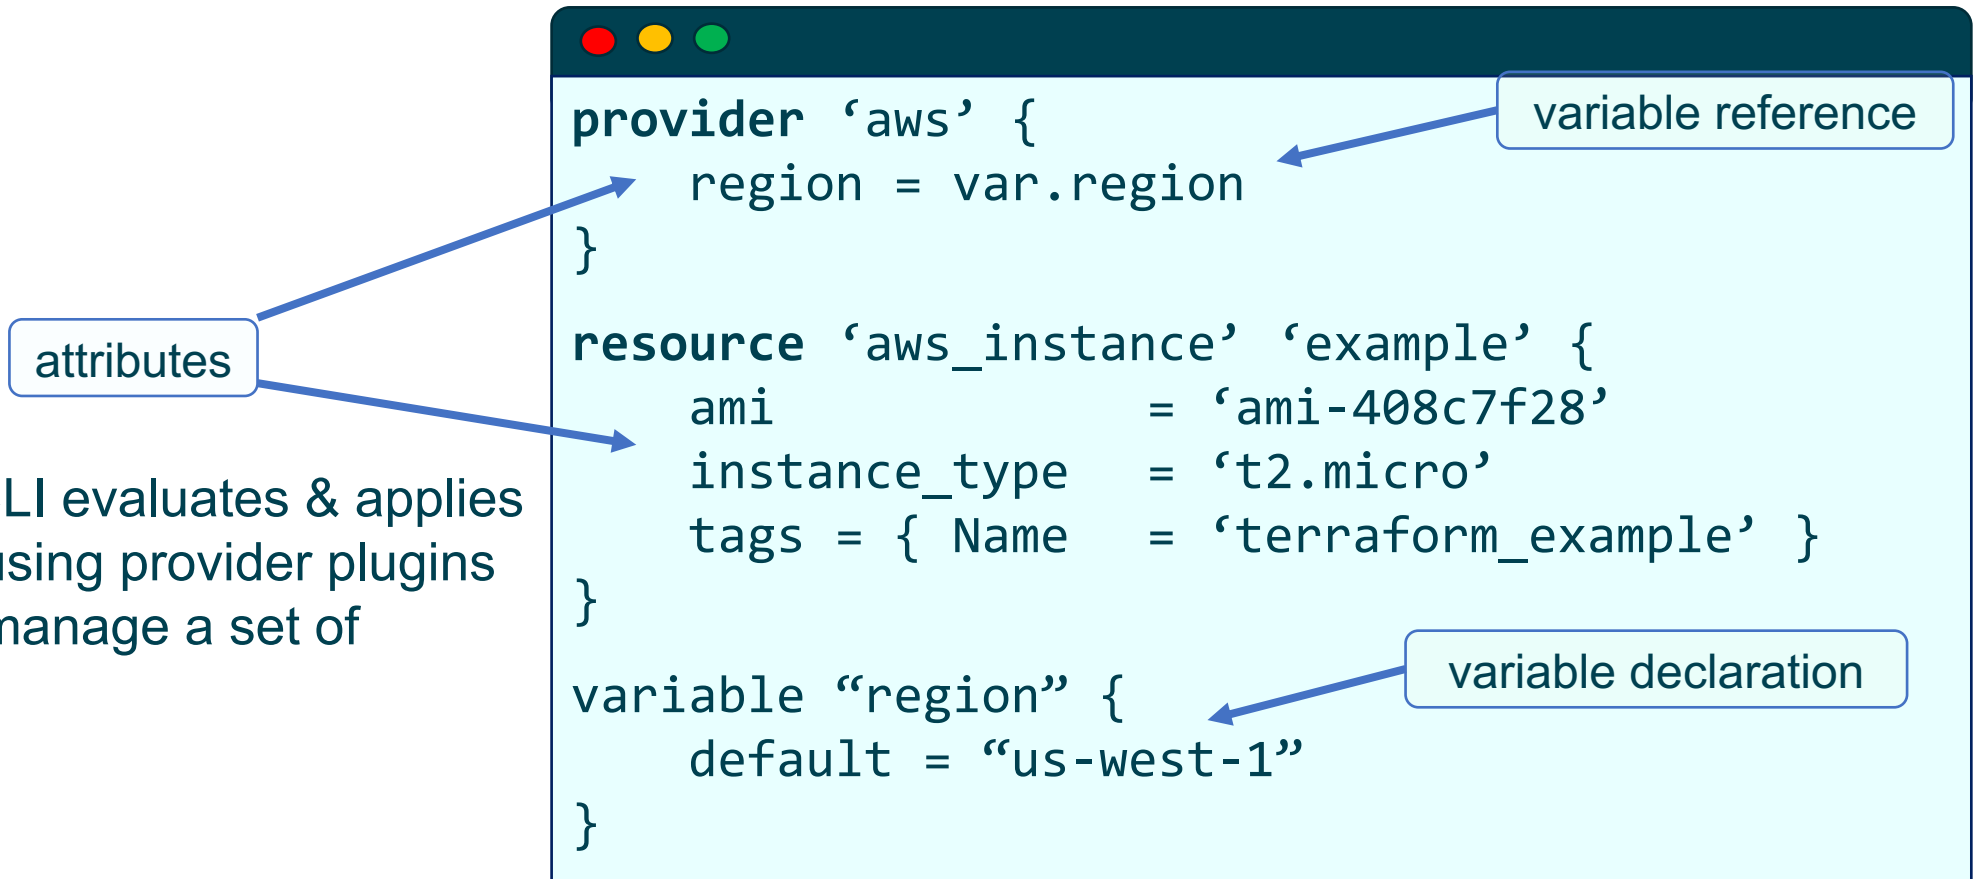
In this model re-applying the same configuration assures the same result, i.e. the same set of resources and state, every time.

# Terraform Configurations

# Example Terraform Configurations

Configurations use Hashicorp Configuration Language (HCL v2) & specify a provider & resources specific to the provider

The Terraform CLI evaluates & applies configurations, using provider plugins which define & manage a set of resource types



JSON format can also be used – files named as `.tf.json`



# (Some) Terraform Providers

## Cloud

**AWS**, Alibaba Cloud, **Azure**, DigitalOcean, Exoscale, **Google Cloud**, Heroku, IBM Cloud, Oracle Cloud, OVH, Packet, 1&1, Spotinst, Linode

## HashiCorp

Vault, Nomad, Consul

## Version Control

GitHub, GitLab, Bitbucket, Azure DevOps

## Orchestrators

**Kubernetes**, Helm, **Docker**, **OpenStack**

## Databases

PostgreSQL, MySQL, MongoDB Atlas

## Hypervisors

VMware NSX-T, vCloud , vSphere, **Proxmox**

## Misc

Cobbler, Datadog, DNS, HTTP, **Local**, **TLS**

# What about you ?

A quick anonymous survey

What Platforms do you currently use ?

What is your experience with IaC ?

What Providers would you like to use with Terraform ?



<https://forms.gle/A7W1DE5or65fJUH59>

# Terraform Installation



# Terraform - Installation

Terraform version

Provider plugins installed – if present  
(installed at '*terraform init*')



A terminal window with a dark blue header bar containing three colored window control buttons (red, yellow, green). The terminal text is as follows:

```
$ terraform version
Terraform v1.2.6
+ provider registry.terraform.io/hashicorp/aws v4.2.0

Your version of Terraform is out of date! Latest version is 1.3.7
Update by downloading at https://www.terraform.io/downloads.html
```


Two blue arrows point from the text boxes above to the terminal output. One arrow points from 'Terraform version' to 'Terraform v1.2.6'. The other arrow points from 'Provider plugins installed...' to '+ provider registry.terraform.io/hashicorp/aws v4.2.0'.

Download latest from: <https://www.terraform.io/downloads.html>

Older or beta releases from: <https://releases.hashicorp.com/terraform/>

@mjbright

# Terraform Workflow



Planning, applying, updating, destroying



<https://linkedin.com/in/mjbright>



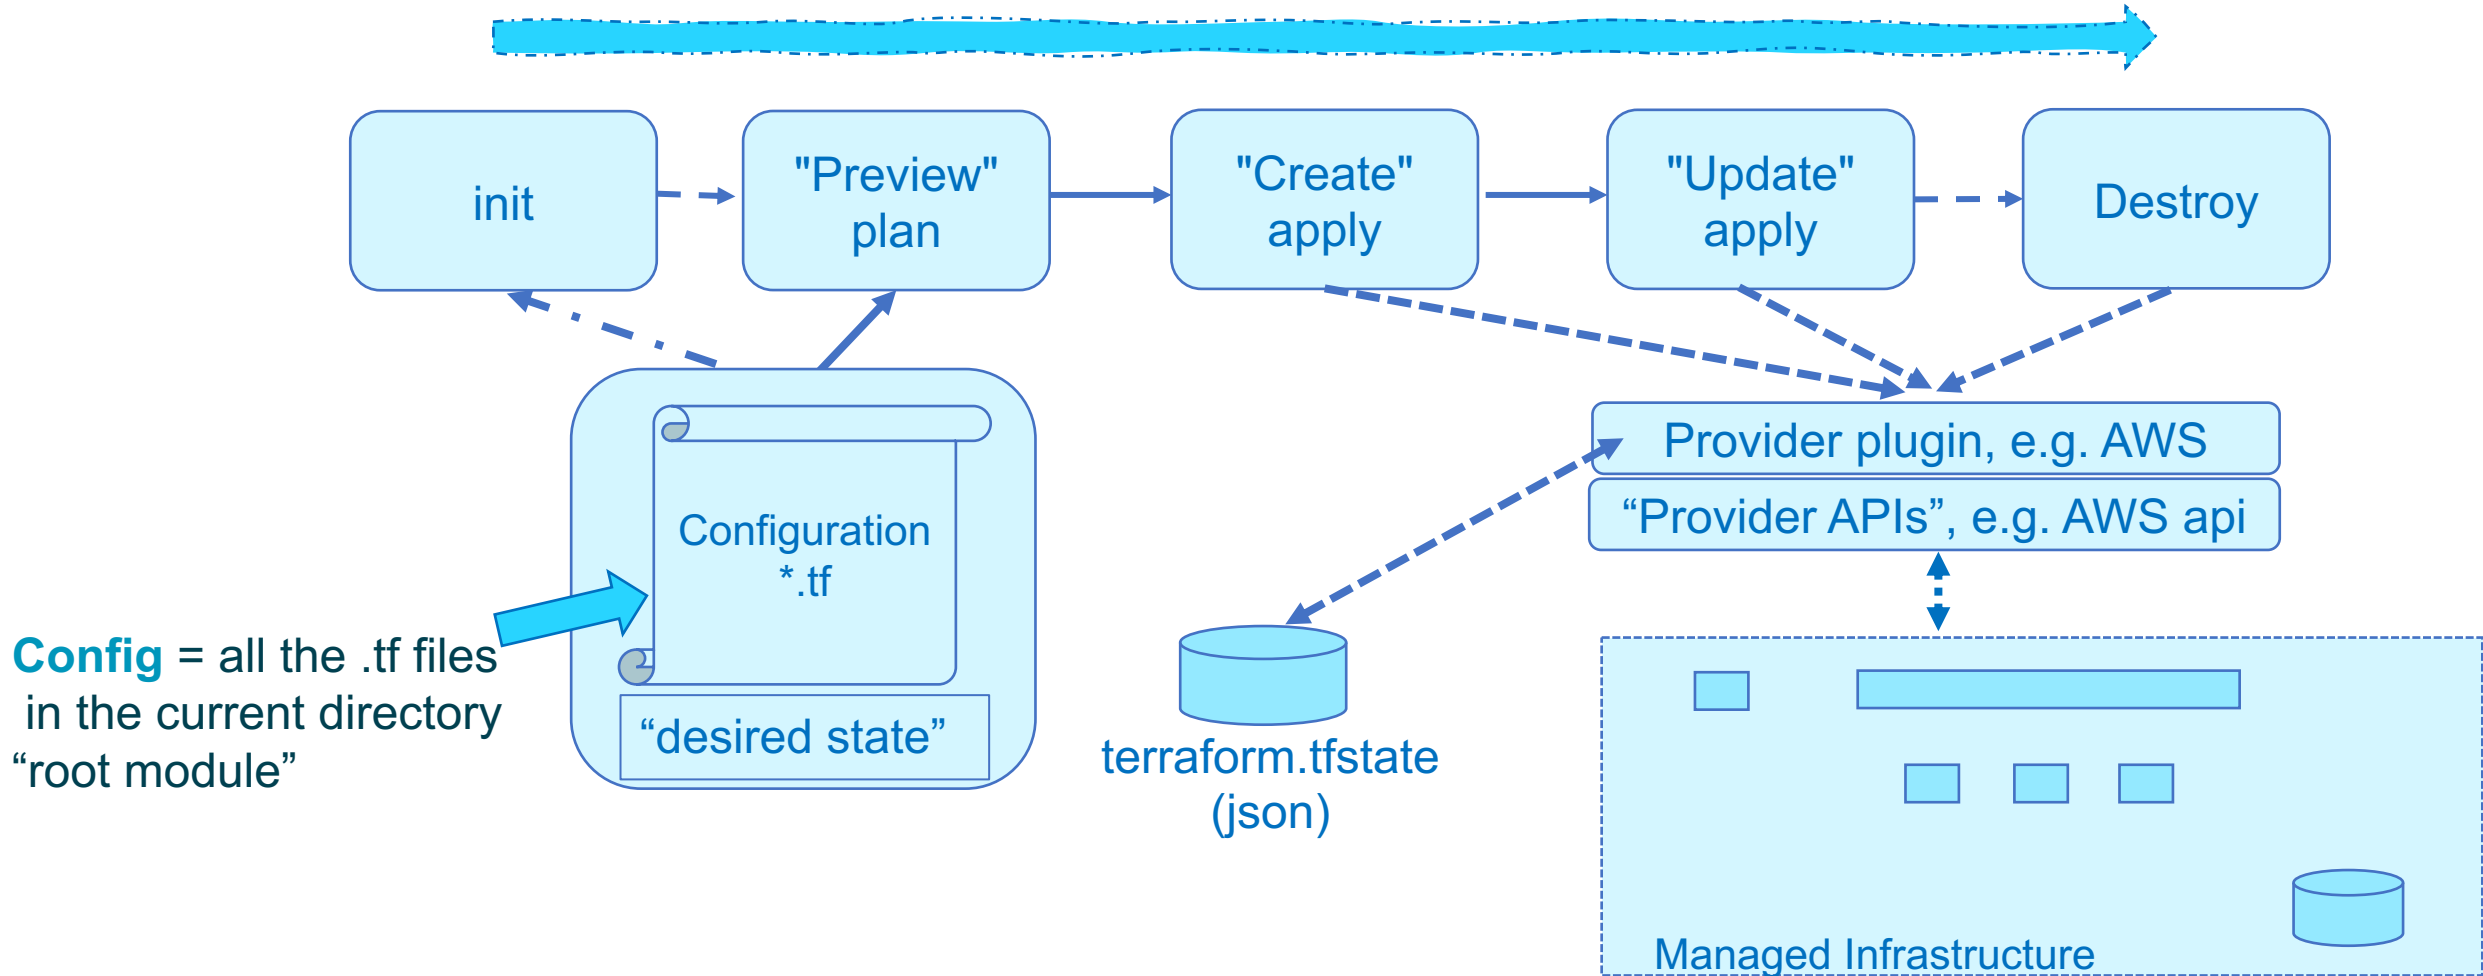
@mjbright



@mjbright

# Terraform Workflow

Terraform determines the current state of the resources it is managing



# What will we do ?



Building up a small 2-tier architecture



<https://linkedin.com/in/mjbright>




@mjbright



@mjbright

# Terraform Workflow



Only a simple architecture right ?

But how about if we can fire it up with 2 commands, in less than 10 seconds ?



# HCL Configs, variables & Control Structures



Parameterizing Terraform configurations through variables



<https://linkedin.com/in/mjbright>



@mjbright



@mjbright

# Terraform Configurations

# Parameterizing Terraform Configurations

We can parameterize our templates using *variables*.

Description, default and type attributes are optional.

```
variable "instance_name" {  
    description = "The name of the EC2 instance"  
}
```

It is best practice is to

Define the variables (as above) in a **variables.tf** file

Specify the default values in a **terraform.tfvars** file

```
instance_name = "resource-value"
```

# Parameterizing Terraform Configurations

Note: the use of the “\${}” syntax to interpolate var.name

```
resource "aws_instance" "example" {  
    ami          = "ami-408c7f28"  
    instance_type = "t2.micro"  
  
    tags = {  
        name = var.name  
        info = "${var.name} info" # string interpolation  
    }  
}
```

# Outputting values

## Outputting values – single instance

We can specify output values in our configuration templates:

```
output "public_ip" { value = aws_instance.ex.public_ip      }
output "public_dns" { value = aws_instance.ex.public_dns    }
output "hosts"      { value = "IP: ${aws_instance.ex.public_ip} Host:
                        ${aws_instance.example.public_dns}" }
```

which will be shown at the end of the apply

Apply complete! Resources: 2 added, 0 changed, 0 destroyed.

Outputs:

public\_dns = ec2-54-183-131-114.us-west-1.compute.amazonaws.com

public\_ip = 54.183.131.114

Hosts = IP: 54.183.131.114 Host: ec2-54-183-131-114.us-west-1.compute.amazonaws.com

and thereafter can be viewed with commands  
*'terraform show'* or *'terraform output'*

# Terraform Types

# Terraform Types

## Types

Terraform supports different types of variables

- Number, String, Boolean
- Lists, Tuples, Sets, Objects
- Maps
- Any to wildcard collections e.g. map(any), list(any)
- null (for omitted values)

<https://www.terraform.io/docs/configuration/types.html>

<https://www.terraform.io/docs/configuration/expressions.html>

# Local Variables



# Locals

Locals are pre-computed values which can be reused elsewhere in the same module

Locals simplify configurations by declaring formulas only once, improving readability if meaningful names are used

```
locals {  
    info_val = "just a reusable string"  
    size     = 10  
    area     = local.size * local.size * 3.1415926835  
}
```

Locals apply across all your configuration files of the current module

```
output test_local { value = local.info_val }  
output circle_area { value = local.area }  
output approx_area { value = "approx. size=${floor(local.area) cm^2}" }
```

# Terraform Control Structures



Abstractions for (looping) creating multiple resources

# Terraform Control Structures

Terraform provides a bewildering array of control structures

- Count loops
- Ternary operator (if)
- For\_each loops
- “for in” loops
- Dynamic blocks
- String templates

# Count Loops

# Terraform Control Structures: “count” loops

Use of the *count* attribute causes creation of multiple resource instances:

```
count = length(some_list)
```

We can access the current index in the resources with `count.index`

Note: elements begin with index zero, not one!

```
resource "type" "name" {  
    count = length(var.listvar)  
  
    resource definition  
    attribute = listvar[ count.index ]  
}
```

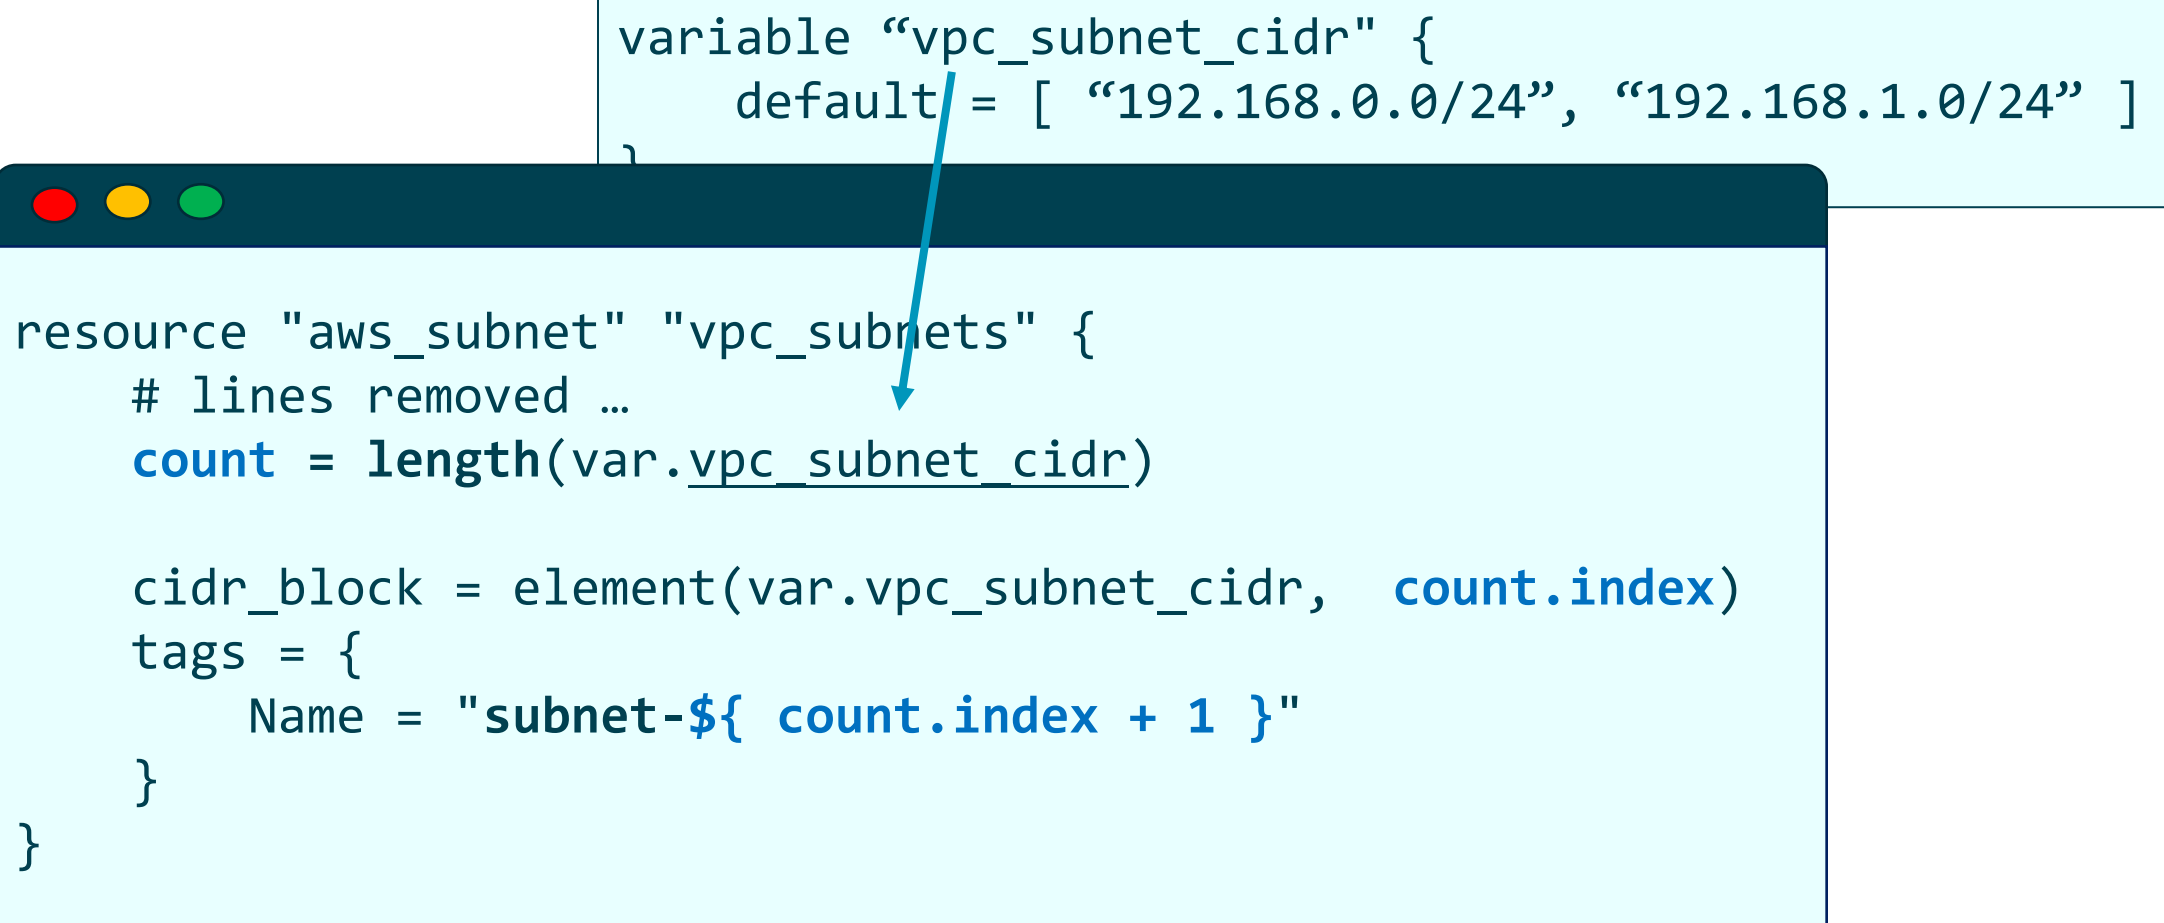
} - count instances created

The created resources can be referenced as `type.name[0]`, `[1]` etc ...

# Terraform Control Structures: “count” loops

Example: To create several aws\_subnet instances (using count)

```
variable "vpc_subnet_cidr" {  
    default = [ "192.168.0.0/24", "192.168.1.0/24" ]  
}
```



```
resource "aws_subnet" "vpc_subnets" {  
    # lines removed ...  
    count = length(var.vpc_subnet_cidr)  
  
    cidr_block = element(var.vpc_subnet_cidr, count.index)  
    tags = {  
        Name = "subnet-${ count.index + 1 }"  
    }  
}
```

**Note:** use of **count.index** to identify the loop number (starts at 0)

@mjbright

# Webinar Agenda

## 1. Introduction to Terraform

- Configurations, Providers
- Installation, Workflow
- Variables: Input, Output, Local

## 2. Control structures

- Replicas using Count

## 3. Working with Terraform Registry

- Data sources
- Modules

## 4. Terraform in Production

- Provisioners
- Managing State

# Data Sources



Querying the Provider for information



# Terraform Data Sources

## Filtering the data

Example: Get the latest Ubuntu 20.04 LTS (for this region)

```
# Get latest Ubuntu Fossa 20.04 AMI
data "aws_ami" "ubuntu-linux-2004" {
  most_recent = true
  owners      = ["099720109477"] # Canonical
  filter {
    name      = "name"
    values    = ["ubuntu/images/hvm-ssd/ubuntu-fossa-20.04-amd64-server-*"]
  }
  filter {
    name      = "virtualization-type"
    values    = ["hvm"]
  }
}
```

# Terraform Data Sources

Example: Using the `aws_ami` data source

```
resource "aws_instance" "server" {  
  name      = "tf-data-source-example"  
  
  key_pair = "terraform"  
  ami      = data.aws_ami.ubuntu-linux-2004.id  
}
```

# Terraform Modules




Modules similar to functions in other languages

# Terraform Modules

## Local source Modules

By convention we define 3 specific terraform files in a module

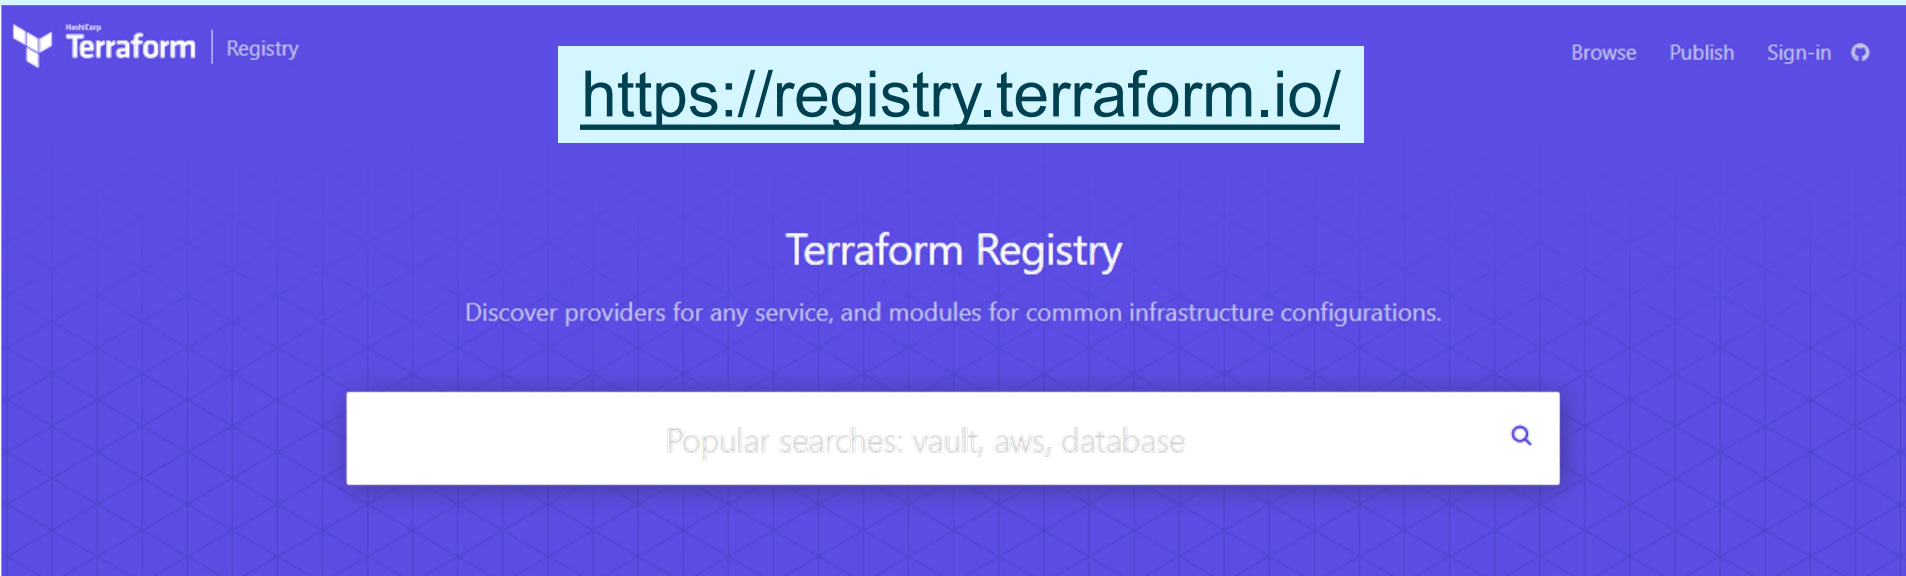


A module may also define its' own local variables

But this is mere convention, we can call these files as we wish.

As modules become more complex it makes sense to split the configuration accordingly


# Terraform Modules - Terraform Registry



## Find Terraform Modules

Use and learn from verified and community modules

### POPULAR MODULES



@mjbright

# Terraform in Production



Provisioners

Importing “foreign” resources

Managing State

Delegating state to remote storage - teams

# Provisioners



Not the same as Providers !

# Terraform Provisioners & user\_data

## Provisioners & user\_data

### User\_data

A common mechanism used to provision cloud resources is **user\_data** which can be used to execute shell commands using Linux **cloud-init** capabilities.

### Provisioners

Another mechanism specific to Terraform is **Provisioners**:

**Local-exec:** Execute code on local machine, which invokes Terraform

**Remote-exec:** Execute code on remote machine managed by Terraform

**File:** Transfer files to the remote machine



# *“Ultimate Terraform”* training Agenda - \*

(for 4 half-day training, Apr 24th to 27th)

## 1. Introduction to Terraform

- Configurations, Providers
- Installation, Workflow
- Variables, types

## 4. Managing State

- Provisioners
- Managing Remote State
- Importing Foreign Resources

## 2. Control structures

- Replicas, templates, ...
- Resource dependencies
- HCL functions, Debugging

## 5. Terraform in Production

- Best Practices
- 3<sup>rd</sup>-party tools
- Terraform Cloud & Enterprise

## 3. Working with Terraform Registry

- Providers & Data sources
- Using & Writing Modules

# Thank you !



<https://linkedin.com/in/mjbright>



@mjbright



@mjbright